

# Linear Mixed-Effects Models (aka Statistics III)

Bernd Figner  
b.figner@psych.ru.nl

Week 5: March 10, 2014

1

## Today: Digging Deeper

- Questions: General, homework-related
- Some homework-related add-ons and leftovers from last class
- $P$  values; multiple cores; post-hocs; writing up model and results; some plotting; linear and quadratic predictors; centering/scaling to reduce multicollinearity
- Advice from Barr et al and others
  - The R-sig-mixed-models debate
  - Testing parameters vs. effects: Bootstrap vs. LRT?
- Homework

2

# **Question 1**

## **lab computers: R or RStudio?**

3

# **Question 2**

## **YOUR data: who's willing to present?**

**March 24 ideal, but other dates possible**

4

# Homework

**Questions? Problems? Comments?**

5

## Some Comments

**In the valuation ratings analysis:**

**Item as a random intercept or not? Why (not)?**

**More theoretical answer**

- They are NOT random sample from all possible items
- They are carefully chosen and created by the experimenter ("repeatable")

**More pragmatic answer**

- We want to test the significance of money and time
- Adding the same predictor as fixed slope and random intercept is typically not a good idea

6

## More Comments

### Time and Money: What should they be?


- Continuous?
- Unordered factors?
- Ordered factors?
- Linear?
- Linear and Quadratic?

7

## Time

### Occurring values

- 0 (= today)
- 3 days
- 14 days
- 28 days
- 42 days



8

## As unordered Factor?

- Assumes that there is no inherent order among the levels
- Seems odd, ignores a lot of information (Money: LOTS of levels, due to "jitter")

For most participants, there will be at least an ordinal relationship:

$$0 < 3 < 14 < 28 < 42$$

9

## As ordered factor?

Recodes the actual values as numeric levels

- today = 1
- 3 days = 2
- 14 days = 3
- 28 days = 4
- 42 days = 5
- polynomial contrasts then test for linear, quadratic (cubic, ...) effects
- More fine-grained information (days): thrown away
- Ordered levels treated as continuous, assuming **equal spacing** between levels (0 vs. 3 = 3 vs. 14??)
- Good: captures linear and higher-order effects: quadratic effects will probably capture unequal spacing

10

## As continuous predictor?

Keeps the fine grained information and unequal spacing between occurring values:

- E.g., difference between 0 and 3 days is smaller than difference between 3 days and 14 days


### To capture non-linear effects

- You need to create a higher-order predictor yourself and add it to the model

11

## Non-Linearity

- **Time: Sooner** → more attractive
- **More money** → more attractive
- **Psychophysics?** Weber-Fechner



12

## Non-Linearity

- **Time: Sooner → more attractive**
- **More money → more attractive**
- **Psychophysics?** Weber-Fechner
  - Effects might not be strictly linear
    - Time: e.g., 0 days vs. 14 days larger than 14 days vs. 28 days?
    - Money: €20 vs €50 larger than €50 vs €80?

**How to capture such possible quadratic effects with continuous predictors?**

13

## Linear and Quadratic IVs

**For the Time IV (same for Money)**

**(1) Center the predictor (or scale)**

```
v5$c_time <- v5$time - mean(v5$time)
```

**(2) Compute squared (quadratic) predictor**

```
v5$q_time <- v5$c_time ^ 2
```

**Do NOT center again**

**(3) Include both in your model**

(don't forget the random slopes for both the linear and quadratic term)

14

## Centering (scaling) reduces collinearity

- Between linear and quadratic variants of the same continuous predictor
- Between main effects and interactions

### Time example

```
v5$q_time_NonCentered <- v5$time ^ 2
```

**Correlation: .964!!!**

```
with(v5, rcor.test(cbind(time, q_time_NonCentered)))
```

	time	q_time_NonCentered
time	*****	<b>0.964</b>
q_time_NonCentered	<0.001	*****

15

## Same with centered predictors

**Correlation: .487**

```
with(v5, rcor.test(cbind(c_time, q_time)))
```

	c_time	q_time
c_time	*****	<b>0.487</b>
q_time	<0.001	*****

**Still not perfect, but much better**

**This is the standard approach**

16



**poly()** creates uncorrelated linear, quadratic, cubic, ... terms

poly(v5\$time, 2) → linear and quadratic terms

```
v5$poly_lin_Time <- poly(v5$time, 2)[,1]
```

```
v5$poly_quadr_Time <- poly(v5$time, 2)[,2]
```

```
with(v5, rcor.test(cbind(poly_lin_Time, poly_quadr_Time)))
```

	poly_lin_Time	poly_quadr_Time
poly_lin_Time	*****	-0.000
poly_quadr_Time	>0.999	*****

**Best approach, as collinearity increases probability of non-convergence (and inflates standard errors, makes interpretation difficult, ...)**

17

- Model with linear and quadratic time and money generated with `poly()`
  - Gives warnings, but produces a model that looks reasonable
- Model with linear and quadratic time and money generated by squaring the centered predictors
  - Needs many more iterations to give reasonably looking model (~ 10,000,000 iterations)
  - More serious warnings even with that many iterations (but results quite similar to poly model)

18

## Model with poly()

```
v_m1c_poly_lin_quad <- lmer(rating ~
poly_lin_Time + poly_quad_Time +
poly_lin_Money + poly_quad_Money +
(1 + poly_lin_Time + poly_quad_Time +
poly_lin_Money + poly_quad_Money | pp_code),
data = v5, control = lmerControl(optCtrl =
list(maxfun = 10000)))
```

### Linear and quadratic predictors for time and money

- In fixed-effects part
- In random part (random slopes)

19

```
summary(v_m1c_poly_lin_quad)
```

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	71.754	2.013	35.65
poly_lin_Time	-176.615	25.538	-6.92
poly_quad_Time	40.375	10.945	3.69
poly_lin_Money	219.521	25.261	8.69
poly_quad_Money	-24.326	10.717	-2.27

### Quadratic term: opposite sign of linear term

- Consistent with Weber-Fechner idea, but better check in plots

20


## Idea: scatter plot with linear and smoothed (quadratic?) trend line

### For time

```
with(v5, plot(time, rating))
with(v5, abline(lm(rating ~ time), col = 'red'))
with(v5, lines(lowess(time, rating), col =
'blue'))
```

- scatter plot with time on x axis and DV on y axis
- adds linear trend line in red
- adds smoothed trend line in blue

21




22

## Same for money

```
with(v5, plot(money, rating))
with(v5, abline(lm(rating ~ money), col = 'red'))
with(v5, lines(lowess(money, rating), col =
'blue'))
```

23



24

**Significance?**

```
Anova(v_m1c_poly_lin_quad, type = 3, test = 'F')
```

	F	Df	Df.res	Pr(>F)	
(Intercept)	1270.9801	1	31	< 2.2e-16	***
poly_lin_Time	47.8298	1	31	0.0000000934307	***
poly_quad_Time	13.6077	1	31	0.0008602	***
poly_lin_Money	75.5162	1	31	0.0000000008191	***
poly_quad_Money	5.1521	1	31	0.0303224	*

25

## Recap and Leftovers from last class

## Significance tests

- **Typically only for fixed effects of interest**
- **Several different options**
- Tests of **coefficients** vs. test of **effects**
- **Coefficients** → like regression: coefficient significantly different from 0?
- **Effects** → like ANOVA: Is whole predictor significant?
  - factor with more than 2 levels
  - whole interaction term
- Post-hoc tests

27

## Tests of Coefficients and „Effects“

- **Example: 3 experimental conditions**
  - neutral
  - positive mood induction
  - negative mood induction
- **Effects:** familiar from ANOVA framework
  - Is the **whole factor** significant?
    - „Effect of mood induction“
  - Similarly: interaction between categorical factors: Is **whole interaction** significant?

28

- **Coefficients:** familiar from regression framework
  - Is the coefficient significantly different from 0?
  - Continuous predictors (→ 1 df)
  - Interactions between continuous predictors (→ 1 df)
  - Factors with 2 levels (→ 1df)
  - **Factors with more levels** (mood induction example)
    - Neutral condition vs average of positive and negative mood significant?
    - Positive vs negative mood condition significant?
    - ...
- **How you set up contrasts matters!**

29

## Getting p values

→ Different methods (usually similar results)

(Douglas Bates: <https://stat.ethz.ch/pipermail/r-help/2006-May/094765.html>)

**List ordered according to recommendations**

<http://glmm.wikidot.com/faq>

30

## Three "most recommended" ones

### (1) (non)Parametric bootstrap

- coefficients: `bootMer()` → `boot.ci()`
- effects: `PBmodcomp()` (package `pbkrtest`)

### (2) Conditional F-tests with df correction

- `Anova(..., test="F")` (package `car`)
- `KRmodcomp()` (package `pbkrtest`)

### (3) Likelihood Ratio Tests

- `anova()` or `drop1()`

31

## Conditional F test with df correction (Kenward-Roger)

```
library(car)
```

```
Anova(mymodel, type = 3, test = "F")
```

To get type 2 test: `type = 2`

`Anova()` calls `KRmodcomp` from `pbkrtest`

```
KRmodcomp(LargeModel, SmallModel)
```

`SmallModel`: same as `LargeModel`, but without the fixed effect of interest

## Notes

- Typically fast; widely accepted
- NOT available for generalized mixed models

32



### Valuation Rating Example

```
Anova(v_m1c, type = 3, test = 'F')
```

Analysis of Deviance Table (Type III Wald F tests with Kenward-Roger df)

Response: rating

	F	Df	Df.res	Pr(>F)
(Intercept)	1270.924	1	31	< 2.2e-16 ***
s_time	47.587	1	31	0.0000000980867 ***
s_money	75.760	1	31	0.0000000007902 ***

### VERY detailed example R script

BlackBoard → Course Documents → Week 4 → Homework Week 4 Example Solution

33

### Bootstrap methods: bootMer()

```
library(lme4)
```

**(1) Run your lmer model (mymodel)**

**(2) Define a function (only once per R session)**

```
FUN_bootMer <- function(fit) {
  return(fixef(fit))
}
```

**(3) Run bootstrap (can take a while!)**

```
boot_mymodel <- bootMer(mymodel, FUN_bootMer,
  nsim = 1000, type = "parametric", .progress =
  "txt", PBargs = list(style = 3))
```

→ Look up ?bootMer

→ Possible to use several cores → much faster!

34

## Bootstrapping can take a long time!

→ Try it first out doing only a few iterations

→ use `Sys.time()` as a stop watch

```
t1 <- Sys.time()
boot_v_m1c <- bootMer(v_m1c, FUN_bootMer, nsim
= 3, type = "parametric")
t2 <- Sys.time()
t2 - t1
```

→ Tells you how long the 3 iterations took

→ Gives you a feel how long 1000 will take...

→ Then run the 1000 iterations (e.g., before going to bed)

35

## Progress bar

```
boot_mymodel <- bootMer(mymodel, FUN_bootMer,
nsim = 1000, type = "parametric", .progress =
"txt", PBargs = list(style = 3))
```

- Shows a simple progress bar plus percentage how many of the simulations R has already finished
- Not possible if you use more than 1 core

36

## Multiple Cores with bootMer()

- detectCores() from library(pbkrttest) → tells you the number of CPUs (cores) of your computer
- bootMer() can use more than 1 → MUCH faster often!

### NOTES

- **Don't use ALL your cores for R!!!**
- **Leave at least 1 for other tasks**

```
boot_mymodel <- bootMer(mymodel, FUN_bootMer,
  nsim = 1000, parallel = "multicore", ncpus =
  3)
```

```
boot_mymodel <- bootMer(mymodel, FUN_bootMer,
  nsim = 1000, parallel = "snow", ncpus = 3)
```

37

## Get Confidence Intervals (for $p$ values)

**Have a look at** boot\_mymodel

```
head(as.data.frame(boot_mymodel))
```

### Get CIs

#### Intercept

```
boot.ci(boot_mymodel, index = 1, conf =
  0.95, type=c("norm", "basic", "perc"))
```

#### 99% CI

```
boot.ci(boot_mymodel, index = 1, conf =
  0.99, type=c("norm", "basic", "perc"))
```

38

**For second column (= first coefficient after intercept)**

```
boot.ci(boot_mymodel, index = 2, conf = 0.95, type=c("norm", "basic", "perc"))
```

**etc for 3<sup>rd</sup> column, 4<sup>th</sup> column, ....**

**When a CI does NOT include 0 → significant!**

- 95% CI does NOT include 0 →  $p < .05$
- 99% CI does NOT include 0 →  $p < .01$
- 95% CI DOES include 0 →  $ns, p > .05$
- 90% CI DOES include 0 →  $ns, p > .10$

39

## Bootstrap method to test EFFECTS

– PBmodcomp(LargeModel, SmallModel) from  
pbkrtest → test of **effects**

**Same idea as KRmodcomp**

- First, fit model with and without effect of interest
- Then, compare the fit of the 2 models via PBmodcomp()

PBmodcomp(LargeModel, SmallModel)

## Notes

- Can take quite some time
- → possible to distribute across several clusters; much faster (check ?PBmodcomp)

40

## Multiple Cores with PBmodcomp()

```
n_cores <- detectCores()
```

**NOTE:** Don't use all cores; leave at least 1 (i.e., `n_cores - 1`)

### Create the clusters

```
clusters <- makeCluster(rep("localhost",  
n_cores - 1))
```

### Run your model comparison

```
PB_1 <- PBmodcomp(mylarge_m, mysmall_m, nsim = 1000, cl  
= clusters)
```

### Look at the result

```
PB_1
```

### Stop the cluster

```
StopCluster(clusters)
```

41

## LRTs

### Again model comparison approach

- First, fit model with and without effect of interest
- Then, compare the fit of the 2 models via `anova()`

```
anova(SmallModel, LargeModel)
```

### Notes

- **Models need to be nested!!**
- LRTs often NOT recommended (for smaller data sets)
- BUT: Barr et al. (2013) → worked very well!
- Recent discussion on R-sig-mixed-models:
- LRTs does not rely on a precise estimate of the coefficients, thus may be better than other methods when model is "overparametrized"

42

- LRTs: models must be fit with ML, not REML
- Fit lmer models with `REML = FALSE`

**BUT: R likes you!**

- R refits the models in the `anova()` command using ML
- Then does model comparison with the ML models

43

## **drop1()**

- If you need to compare many models

```
drop1_mymodel_Chisq <- drop1(mymodel,  
~., test = "Chisq")
```

- Does all the relevant `anova()` model comparisons (i.e., LRTs) for you.

44

## Post-hoc tests

- If you have a factor with more than 2 levels
- **Age Group**
  - Children
  - Adolescents
  - Adults
- **You know how to test whether Age Group as a whole is significant**
- **How to do pairwise post-hoc comparisons?**
  - Children vs. Adolescents
  - Children vs. Adults
  - Adolescents vs. Adults

45

```
myAgeModel <- lmer(DV ~ f_Agegroup + IV + (1 + IV | f_pp), data = mydata)
```

### Simple way to do post-hoc tests for f\_Agegroup

```
library(lsmeans)
posthoc_Agegroup <- lsmeans(myAgeModel,
list(pairwise ~ f_Agegroup))
posthoc_Agegroup
```

### Same, but glht() from multcomp

```
library(multcomp)
posthoc_Agegroup <- glht(myAgeModel, linfct =
mcp(f_Agegroup = "Tukey"))
summary(posthoc_Agegroup)
```

46

### An output example (from one of my studies)

```
$`f_Agegroup lsmeans`
f_Agegroup    lsmean      SE      df lower.CL upper.CL
child 59.42722 1.611361 193.7102 56.24916 62.60529
adolescent 65.61498 1.413526 198.4173 62.82752 68.40245
adult 66.37734 1.613881 194.8105 63.19442 69.56027

$f_Agegroup pairwise differences`
              estimate      SE      df  t.ratio p.value
child - adolescent -6.1877614 2.143488 195.7366 -2.88677 0.01200
child - adult      -6.9501220 2.280591 194.2601 -3.04751 0.00738
adolescent - adult -0.7623606 2.145383 196.3642 -0.35535 0.93279
p values are adjusted using the tukey method for 3 means
```

47

### Post-hocs for factor interactions

```
mymodel <- lmer(DV ~ f_1 * f_2 + ....)
(1) Create one single new factor that combines the
two predictors IV1 and IV2
f_1 <- c('a', 'a', 'b', 'b')
f_2 <- c('d', 'e', 'e', 'd')
int_f_12 <- interaction(f_1, f_2)
[1] a.d a.e b.e b.d
Levels: a.d b.d a.e b.e
same: paste(f_1, f_2, sep = '_')
```

48



**(2) Run the your model again, but use the newly created factor instead of `f_1 * f_2`**  
`mymodel_2 <- lmer(DV ~ int_f_12 + ....)`

**(3) Use `lsmeans()` or `glht()` to do all pairwise post-hoc comparisons**

`lsmeans(mymodel_2, list(pairwise ~ int_f_12))`

`glht(mymodel_2, linfct = mcp(int_f_12 = "Tukey"))`

49

## How to write up a model and the results

### **Mixed-models are relatively new methods** (e.g., „mixed-model“ vs „multilevel“)

- No established standards how to write it up (no APA guidelines for mixed-models)
- Some examples: BlackBoard → Course Documents → Examples of papers using lme4
- Example how I tend to write up my results (made-up)  
→ It's also on BlackBoard

**NOTE:** cite R and and packages you use!

`citation()` → Citation info for R

`citation("lme4")` → Citation info for lme4

`citation("car")` → Citation info for car

...

51

## **Model Setup**

The valuation rating data were analyzed with a linear mixed-effects model approach, using the lmer function of the lme4 package (version 1.1.-4; Bates, Maechler, Bolker, & Walker, 2014) in R (R Core Team, 2013). The model included a fixed intercept and each a fixed slope for the continuous predictors Time and Money. These continuous predictors were scaled before we entered them into the model (i.e., they had a mean of 0 and a standard deviation of 1).

52

**cont.**

We followed Barr, Levy, Scheepers, and Tily's (2013) advice to use a maximal random-effects structure: The repeated-measures nature of the data was accordingly modeled by including a per-participant random adjustment to the fixed intercept ("random intercept"), as well as per-participant random adjustments to the Time and Money slopes ("random slopes"); in addition, we included all possible random correlation terms among the random effects.

53

***P* Values**

*P* values were determined using conditional F tests with Kenward-Roger correction of degrees-of-freedom, as implemented in the Anova function (with Type III F tests) from the package car (version 2.0.19; Fox & Sanford, 2011; this function calls the KRmodcomp function from the package pbkrtest: Halekoh & Højsgaard, 2013).

54

**OR:**  $P$  values were determined using Likelihood Ratio Tests as implemented in the anova function of the R base package.

**OR:**  $P$  values were determined using parametric bootstrapping as implemented in lme4's bootMer function, with 1000 simulations and deriving confidence intervals using the function boot.ci of the package boot (version 3.1.9.; Canty & Ripley, 2013; Davison & Hinkley, 1997).

**OR:** ... whatever you used!

55

## Results

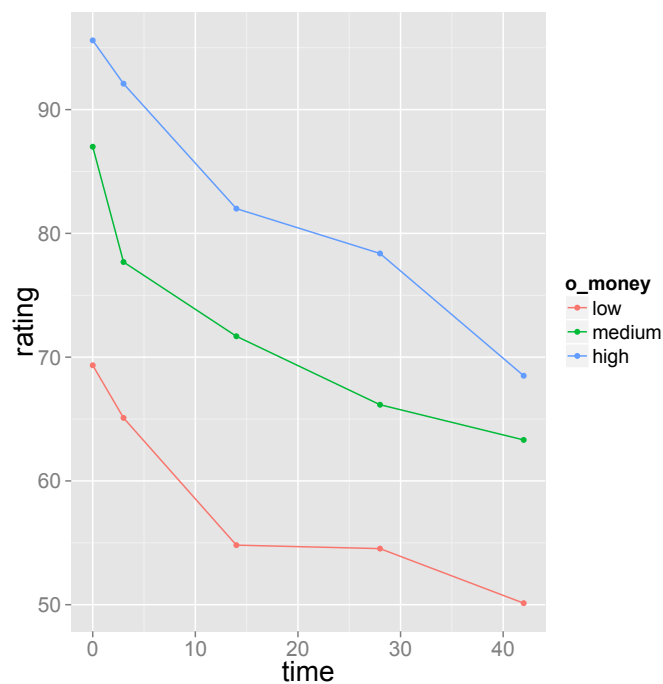
As expected, we found significant effects of Time (coef = -7.965;  $p < .001$ ) and Money (coef = 10.038;  $p < .001$ ): As time-of-delivery increased, the valuation ratings decreased and as Money increased, the valuation ratings increased. Figure 1 shows the effects of Time and Money on the valuation ratings. ETC ETC...

56

## Figures: make them pretty...

```
ggplot(data = v5, aes(x = time, y = rating,
  colour = o_money, group = o_money)) +
  stat_summary(fun.y = mean, geom = "point") +
  stat_summary(fun.y = mean, geom = "line") +
  theme(axis.text.x = element_text(size = 15),
  axis.title.x = element_text(size = 20),
  axis.text.y = element_text(size = 15),
  axis.title.y = element_text(size = 20),
  legend.title = element_text(size = 14),
  legend.text = element_text(size = 14))
```

57



58

## Make them prettier...

- While this is better than my usual ugly ones, there's room for further improvement:
- Categories on the x axis
- Title of the legend (o\_money) and its position
- ...

## Error bars?

- Long and complicated story (mixed designs!)
- I don't add error bars usually
- Recent advice → if you want to show uncertainty, plot the coefficients and **their** error bars

59

## xypplot() for individual participant patterns



### For Money

```
xypplot(rating ~ money | pp_code, groups
= pp_code, data = v5, type = c('p',
'r'))
```

### For Time

```
xypplot(rating ~ time | pp_code, groups =
pp_code, data = v5, type = c('p', 'r'))
```

60




## Time **and** Money

Time on x axis; **o**\_money as separate lines

```
xyplot(rating ~ time | pp_code, groups =  
o_money, data = v5, type = c('p', 'r'),  
auto.key = TRUE)
```

**auto.key = TRUE** → adds figure legend

63



64

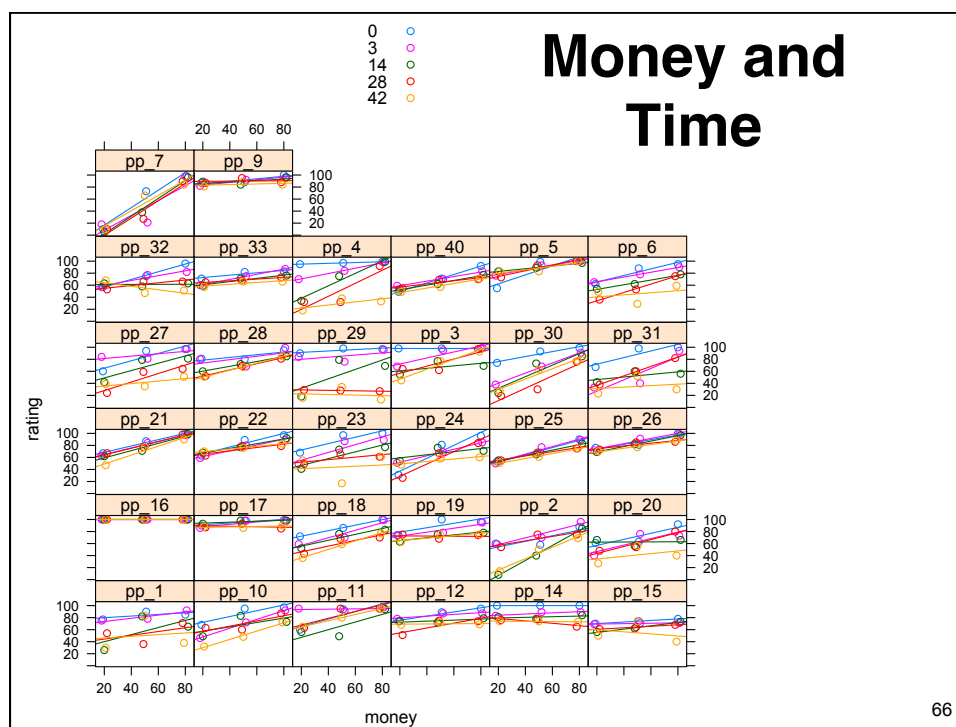


## Money and Time

**Money on x axis; Time as separate lines**

```
xyplot(rating ~ money | pp_code, groups
= time, data = v5, type = c('p', 'r'),
auto.key = TRUE)
```

65



66

## Individuals differ a lot in effects of Time and Money

It would be nice to get the numbers how the participants differ from each other (i.e., the random intercepts and random slopes)

### **fixef()**

→ **fixed effect estimates of your model**

```
> fixef(v_m1c)
(Intercept)      s_time      s_money
  71.754167    -7.964761    10.037750
```

67

### **ranef()**

→ **per-participant adjustments to fixed effects**

```
> ranef(v_m1c)
$pp_code
      (Intercept)      s_time      s_money
pp_1    -5.616019409   -4.7786399  -1.6513253
pp_10   -2.931644297   -1.0844274   2.3557412
pp_11    9.114218849    5.5266371   0.2360635
pp_12    4.772589075    0.6948460  -4.0505320
pp_14   10.897122095   -0.4130446  -8.6088317
...
```

68

**coef()**

→ random and fixed effects together

i.e., best estimate for each pp's coefficients

```
> coef(v_m1c)
$pp_code
      (Intercept)      s_time      s_money
pp_1      66.13815 -12.74340079   8.3864244
pp_10     68.82252  -9.04918823  12.3934909
pp_11     80.86839  -2.43812376  10.2738132
pp_12     76.52676  -7.26991482   5.9872177
pp_14     82.65129  -8.37780546   1.4289180
pp_15     67.73749  -8.60869805   5.7816330
...
```

69

**If you have more than 1 grouping variable**

```
coef(mymodel)$pp_code
```

```
coef(mymodel)$f_item
```

**To save the coefs in a csv file**

```
v_m1c_coefs <- coef(v_m1c)$pp_code
```

```
write.csv(v_m1c_coefs, row.names = TRUE, file
= 'v_m1c_coefs.csv')
```

row.names = TRUE → pp\_code is included as 1<sup>st</sup> column

**Why useful to save them?**

- Descriptives, plots, ...
- Use in some other analyses (e.g., fMRI)

70

	A	B	C	D
1	(Intercept)	s_time	s_money	
2	pp_1	66.1381472572	-12.7434007914	8.3864244114
3	pp_10	68.82252237	-9.0491882322	12.3934908987
4	pp_11	80.8683855158	-2.438123758	10.2738132174
5	pp_12	76.5267557415	-7.2699148168	5.9872177209
6	pp_14	82.6512887612	-8.3778054582	1.4289179751
7	pp_15	67.7374862816	-8.6086980507	5.7816329675
8	pp_16	98.4146427397	0.0584053472	0.1251181152
9	pp_17	91.2446027856	-2.8344483949	2.3280414432
10	pp_18	69.4885297045	-9.0795906928	12.4211233642
11	pp_19	77.5661384166	-6.4604392642	6.533258789
12	pp_2	60.646298243	-8.2973887227	17.0735410335
13	pp_20	61.60410219	-10.4891250255	10.9044332116
14	pp_21	77.8488171908	-3.3717236554	12.5630517301
15	pp_22	77.598557432	-2.9070871406	9.754576843
16	pp_23	65.8868752597	-13.1229341827	10.3932670512
17	pp_24	64.5482595516	-6.5455323517	14.940972429
18	pp_25	68.5583559368	-4.8510239655	12.7445090569
19	pp_26	82.2097265985	-3.1033641951	8.3688932814
20	pp_27	65.1912963837	-15.4904477586	10.5387662904
21	pp_28	74.7331485863	-6.7622635605	9.4681826219
22	pp_29	57.07762275	-25.2000636475	6.3157032038

71

## Today: Digging Deeper

- Questions: General, homework-related
- Some homework-related add-ons
- Leftovers from last class: p values; multiple cores
- Post-hocs; writing up model and results; some plotting; linear and quadratic predictors; centering/scaling to reduce multicollinearity
- **Advice from Barr et al and others**
  - The R-sig-mixed-models debate
  - Testing parameters vs. effects: Bootstrap vs. LRT?
- Homework

72

## R-sig-mixed models

### Who has read these emails recently?

- If not, do it as homework (available online as archives)  
→ debate started on Feb 28, 2014

### Short summary

- **E. Ellsiepen:** Followed Barr et al advice: maximal random-effects structure + LRTs for p values; odd results
- **D. Bates:** model is overparametrized (too complex for data); "Barr et al advice is dangerous"
- **R. Levy** (from Barr et al): perhaps no overparametrization; Bates misunderstood advice in paper; LRTs are fine even when overparametrized (testing significance of coefficient more problematic!)
- **D. Bates:** I should have re-read the paper

73

- **R. Levy:** More detailed advice about strategies; suggestions: try different optimizer; use LRTs, not tests that rely on precise estimate of coefficient
- Maechler, Bolker, ... thread is still alive

### Important thoughts

#### If model overparametrized/with convergence problems:

- Try different optimizer(s): glmer: nlminb from package optimx  
`optimizer="optimx", optCtrl=list(method="nlminb")`
- **Use LRTs for p values** → For model comparison approach, not important whether the really-really best value for a coefficient is found!  
→ Different coefficient values give virtually equally good fit  
→ Actual value of coeff irrelevant for comparison of models with/without the predictor → more trustworthy p values!

74